

### Assignment -3

Assignment Date	7 October 2022
Student Name	Sivaleshwari.M
Student Roll Number	812419106045
Maximum Marks	2 Marks

```
In [121]: import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

from sklearn.linear_model import LinearRegression
from sklearn import metrics

%matplotlib inline
```

```
In [122]: os.getcwd()
```

```
Out[122]: 'C:\\Users\\pc'
```

```
In [123]: path='C:\\Users\\pc\\downloads\\'
data=pd.read_csv(path+'abalone.csv')
data
```

```
Out[123]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7

ID	Sex	Length	Depth	Width	Height	Weight	Length2	Length3	Age
4172	F	0.565	0.150	0.165	0.0079	0.3700	0.2200	0.2180	11
4173	M	0.590	0.440	0.135	0.0980	0.4350	0.2145	0.2805	10
4174	M	0.600	0.475	0.205	1.1760	0.6285	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	M	0.710	0.525	0.195	1.9485	0.8455	0.3765	0.4050	12

4177 rows x 10 columns

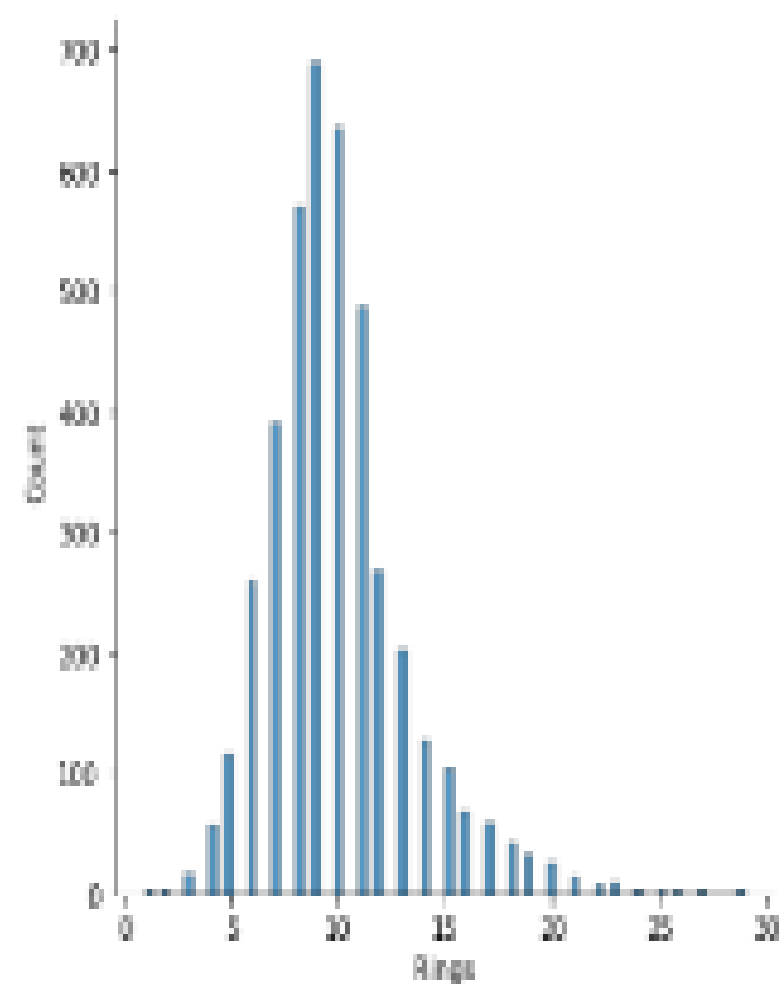
```
In [124]: data.shape
```

```
Out[124]: (4177, 10)
```

```
In [125]: sns.displot(data['Rings'])
```

```
Out[125]: cassborn.axisgrid.FacetGrid at 0x2bef051b370
```

```
Out[125]: cassborn.axisgrid.FacetGrid at 0x2bef051b370
```

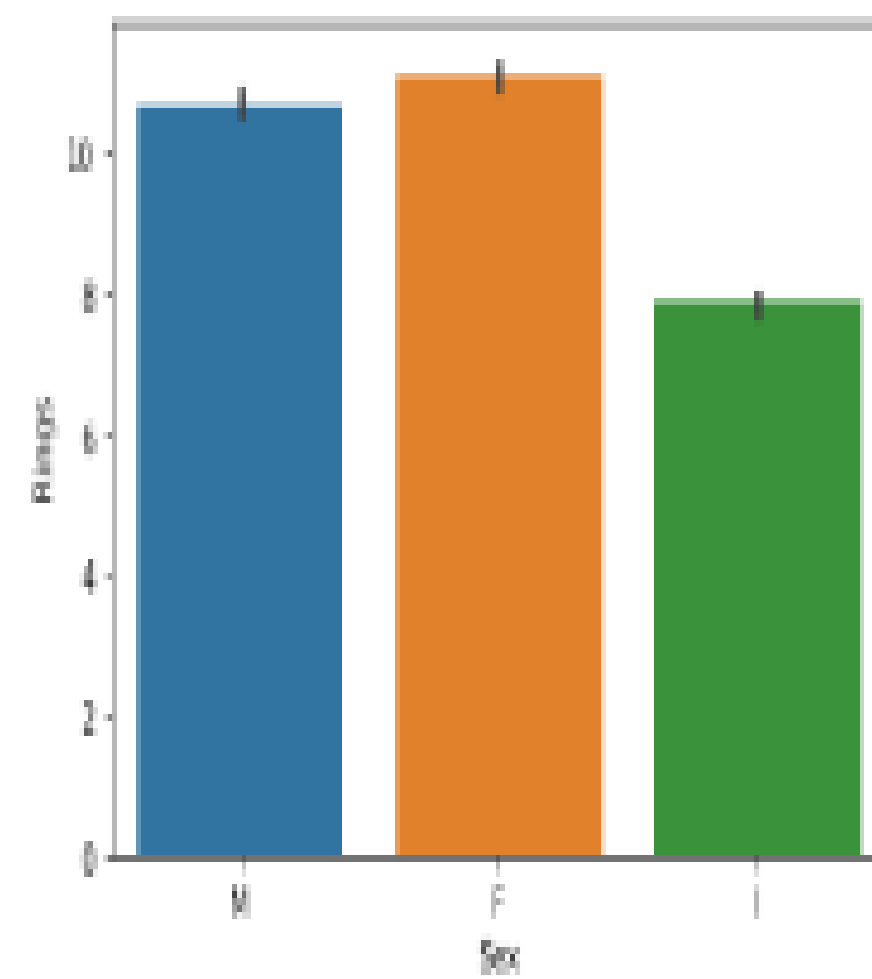


```
In [126]: sns.barplot(x='Sex', y='Rings', data=data)
```

```
Out[126]: AxesSubplot: xlabel='Sex', ylabel='Rings'
```

```
In [126]: sns.barplot(x='Sex',y='Rings',data=data)
```

```
Out[126]: <AxesSubplot:xlabel='Sex', ylabel='Rings'>
```



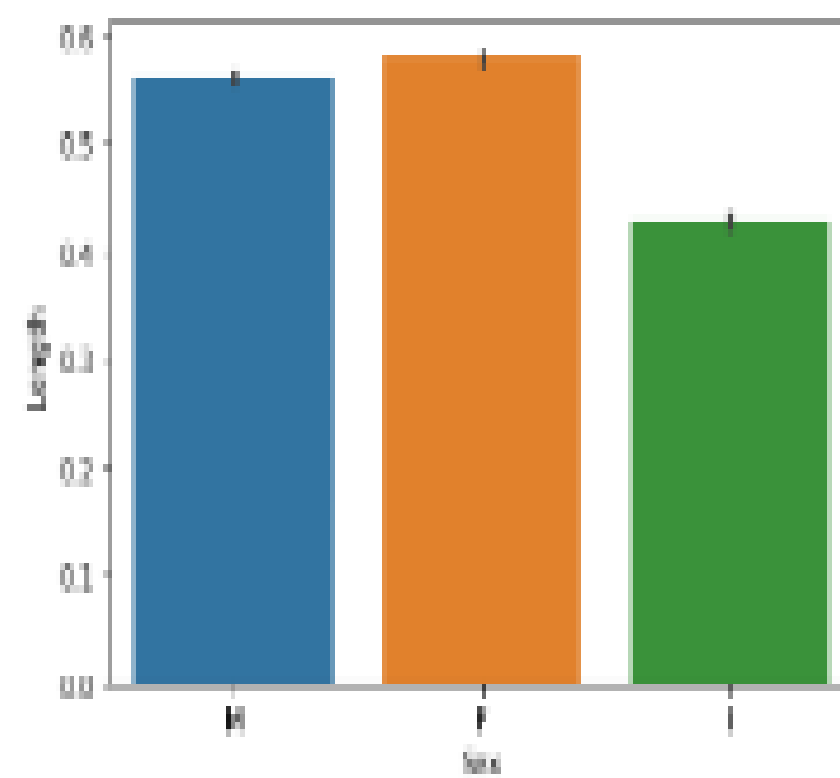
```
In [127]: sns.barplot(x='Sex',y='length',data=data)
```

```
Out[127]: <AxesSubplot:xlabel='Sex', ylabel='length'>
```

...

```
In [128]: sns.barplot(x='Sex',y='length',data=data)
```

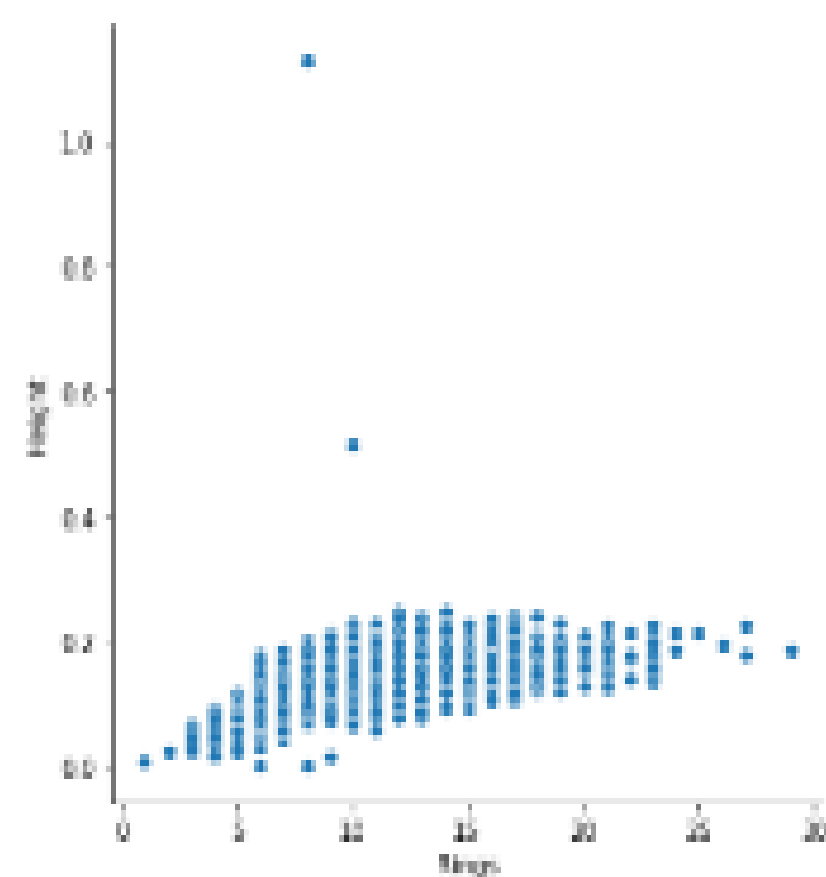
```
Out[128]: <AxesSubplot:xlabel='Sex', ylabel='length'>
```



```
In [129]: sns.relplot(x='Rings',y='height',data=data)
```

```
Out[129]: <seaborn.axisgrid.FacetGrid at 0x2be78a8880>
```

Out[128]: <seaborn.axisgrid.FacetGrid at 0x2bef70a8000>



```
In [129]: import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)# to avoid warning
```

```
In [129]: import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)# to avoid warning
```

```
In [130]: sns.pairplot(data=data[['sex', 'Length', 'Diameter', 'height', 'whole weight', 'kings']], hue='kings')
```

```
c:\users\pc\anaconda\lib\site-packages\seaborn\distributions.py:150: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
c:\users\pc\anaconda\lib\site-packages\seaborn\distributions.py:150: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
c:\users\pc\anaconda\lib\site-packages\seaborn\distributions.py:150: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
c:\users\pc\anaconda\lib\site-packages\seaborn\distributions.py:150: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
c:\users\pc\anaconda\lib\site-packages\seaborn\distributions.py:150: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
c:\users\pc\anaconda\lib\site-packages\seaborn\distributions.py:150: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
c:\users\pc\anaconda\lib\site-packages\seaborn\distributions.py:150: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
```

```
In [130]: data.isnull().sum()
```

```
Out[130]: sex          0
Length          0
Diameter        0
height          0
```

```
Out[130]: Sex          0
Length          0
Diameter        0
Height          0
Whole weight    0
Shucked weight  0
Viscera weight  0
Shell weight    0
Rings           0
dtype: int64
```

```
In [131]: data.describe()
```

```
Out[131]:
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523982	0.407601	0.128516	0.828742	0.356367	0.100584	0.216631	9.930604
std	0.120063	0.080290	0.141827	0.490388	0.221963	0.088614	0.136203	3.224189
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
25%	0.450000	0.350000	0.015000	0.441500	0.150000	0.003500	0.130000	5.000000
50%	0.545000	0.425000	0.140000	0.798500	0.330000	0.171000	0.234000	9.000000
75%	0.615000	0.480000	0.185000	1.153000	0.502000	0.253000	0.325000	11.000000
max	0.815000	0.650000	1.130000	2.025500	1.400000	0.760000	1.005000	29.000000

```
In [132]: data.skew()
```

```
Out[132]: Length          -0.639873
Diameter          -0.609198
Height            3.128817
Whole weight      0.510659
Shucked weight    0.719698
Viscera weight    0.591852
Shell weight      0.628927
Rings             1.114183
dtype: float64
```

```
In [133]: sns.boxplot(data['Rings'])
```

```
TypeError                                Traceback (most recent call last)
<ipython-input-133-8f1dd4be65a1> in <module>
----> 1 sns.boxplot(data['Rings'])

TypeError: list indices must be integers or slices, not str
```

```
In [134]: q1 = data['Rings'].describe()['25%']
          q3 = data['Rings'].describe()['75%']
          qI
```

```
Out[134]: 8.0
```

Type Markdown and LaTeX:  $u^2$

```
In [84]: u2
```

```
Out[84]: 11.0
```

```
In [85]: iqr=q3-q1
iqr
```

Out[85]: 3.8

```
In [86]: a=q1-(1.5*q1)
b=q3+(1.5*q3)
print(a, b)
```

-4.8 37.5

```
In [87]: data[data['Rings']<a]
```

Out[87]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
--	-----	--------	----------	--------	--------------	----------------	----------------	--------------	-------

```
In [88]: data[data['Rings']>b].head()
```

Out[88]:

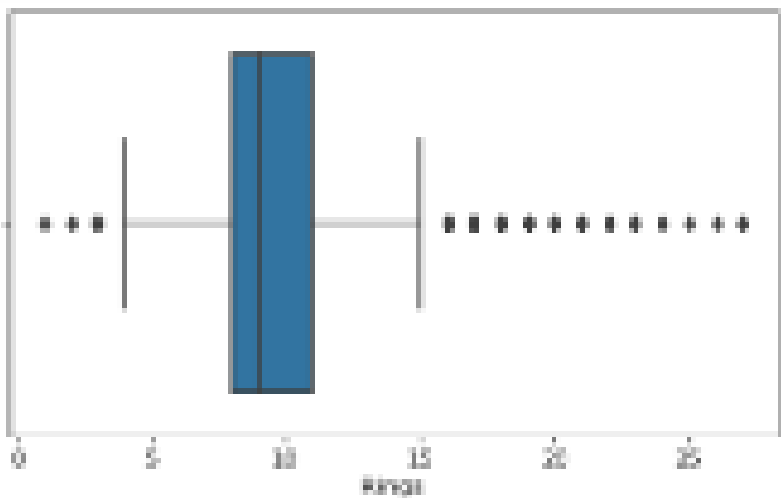
	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
400	F	9.7	8.585	0.105	1.0875	0.7035	0.3215	0.475	28

```
In [89]: outlier_list=list(data[data['Rings']>b]['Rings'])
print(outlier_list)
```

[29]

```
In [90]: data['Rings']=data['Rings'].replace(outlier_list)
sns.boxplot(data['Rings'])
```

Out[90]: <AxesSubplot: xlabel='Rings'>



```
In [90]: le=LabelEncoder()
data['Sex']=le.fit_transform(data['Sex'])
data
```

Out[90]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1530	15
1	2	0.350	0.295	0.090	0.2255	0.0995	0.0905	0.0710	7
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	2	0.440	0.385	0.105	0.5180	0.2185	0.1140	0.1650	16
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...									
4172	0	0.505	0.450	0.165	0.6870	0.3780	0.2300	0.2400	11
4173	2	0.490	0.440	0.155	0.6680	0.4550	0.2145	0.2635	16

4174	2	0.000	0.475	0.205	1.1768	0.5255	0.2875	0.3088	9
4176	0	0.025	0.405	0.158	1.0915	0.5318	0.3610	0.2968	10
4178	2	0.710	0.555	0.195	1.0485	0.9455	0.3785	0.4958	12

4177 rows x 9 columns

```
In [96]: x = data.iloc[:, :-1].values
y = data.iloc[:, -1].values
```

```
In [97]: train_X, val_X, train_y, val_y = train_test_split(X, y, test_size = 0.2, random_state = 0)
print("Shape of training X :", train_X.shape)
print("Shape of Validation X :", val_X.shape)

Shape of Training X : (3341, 8)
Shape of Validation X : (836, 8)
```

```
In [98]: lr = LinearRegression()
lr.fit(train_X, train_y)
print('Attempting to fit Linear Regressor')
```

Attempting to fit Linear Regressor

```
In [100]: %%time
y_pred_val_lr = lr.predict(val_X)
print('MAE on Validation set :', metrics.mean_absolute_error(val_y, y_pred_val_lr))
print("\n")
print('MSE on Validation set :', metrics.mean_squared_error(val_y, y_pred_val_lr))
print("\n")
print('RMSE on Validation set :', np.sqrt(metrics.mean_squared_error(val_y, y_pred_val_lr)))
print("\n")
print('R2 Score on Validation set :', metrics.r2_score(val_y, y_pred_val_lr))
print('R2 score on Validation set :', metrics.r2_score(val_y, y_pred_val_lr))
print("\n")
```

MAE on Validation set : 1.570004956008962812

RMSE on Validation set : 4.7449590677450635

RMSE on Validation set : 1.2061971064638882

R2 score on Validation set : 0.41661000000000007

Wall time: 8 ms

To | |