

Assignment -3
Python Programming

Assignment Date	7 October 2022
Student Name	Sowmiya T
Student Roll Number	812419106046
Maximum Marks	2 Marks

```
In [121]: import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

from sklearn.linear_model import LinearRegression
from sklearn import metrics

%matplotlib inline
```

```
In [122]: os.getcwd()
```

```
Out[122]: 'C:\\Users\\pc'
```

```
In [123]: path='C:\\Users\\pc\\downloads\\'
data=pd.read_csv(path+'abalone.csv')
data
```

```
Out[123]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7

...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12

4177 rows x 9 columns


```
In [124]: data.shape
```

```
Out[124]: (4177, 9)
```

```
In [125]: sns.displot(data['Rings'])
```

```
Out[125]: <seaborn.axisgrid.FacetGrid at 0x2bef651b370>
```

```
Out[125]: <seaborn.axisgrid.FacetGrid at 0x2bef651b370>
```




```
In [126]: sns.barplot(x='Sex', y='Rings', data=data)
```

```
Out[126]: <AxesSubplot:xlabel='Sex', ylabel='Rings'>
```

```
In [126]: sns.barplot(x='Sex',y='Rings',data=data)
```

```
Out[126]: <AxesSubplot:xlabel='Sex', ylabel='Rings'>
```




```
In [127]: sns.barplot(x='Sex',y='Length',data=data)
```

```
Out[127]: <AxesSubplot:xlabel='Sex', ylabel='Length'>
```

..

```
In [127]: sns.barplot(x='Sex',y='Length',data=data)
```


```
Out[127]: <AxesSubplot:xlabel='Sex', ylabel='Length'>
```



```
In [128]: sns.relplot(x='Rings',y='Height',data=data)
```

```
Out[128]: <seaborn.axisgrid.FacetGrid at 0x2bef76a8880>
```

Out[128]: <seaborn.axisgrid.FacetGrid at 0x2bef76a8880>



```
In [129]: import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)# to avoid warning
```

```
In [129]: import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)# to avoid warning
```

```
In [55]: sns.pairplot(data= data [['Sex', 'Length', 'Diameter', 'Height', 'Whole weight', 'Rings']],hue='Rings')
```

```
C:\Users\pc\anaconda3\lib\site-packages\seaborn\distributions.py:305: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
C:\Users\pc\anaconda3\lib\site-packages\seaborn\distributions.py:305: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
C:\Users\pc\anaconda3\lib\site-packages\seaborn\distributions.py:305: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
C:\Users\pc\anaconda3\lib\site-packages\seaborn\distributions.py:305: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
C:\Users\pc\anaconda3\lib\site-packages\seaborn\distributions.py:305: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
C:\Users\pc\anaconda3\lib\site-packages\seaborn\distributions.py:305: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
C:\Users\pc\anaconda3\lib\site-packages\seaborn\distributions.py:305: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
C:\Users\pc\anaconda3\lib\site-packages\seaborn\distributions.py:305: UserWarning: Dataset has 0 variance; skipping density estimate.
  warnings.warn(msg, UserWarning)
```

```
In [130]: data.isnull().sum()
```

```
Out[130]: Sex          0
Length        0
Diameter      0
Height        0
```

```
Out[130]: Sex          0
Length          0
Diameter        0
Height          0
Whole weight    0
Shucked weight  0
Viscera weight  0
Shell weight    0
Rings           0
dtype: int64
```

```
In [131]: data.describe()
```

```
Out[131]:
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	9.933684
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	8.000000
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	9.000000
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	11.000000
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	29.000000

```
In [132]: data.skew()
```

```
Out[132]: Length      -0.639873
Diameter      -0.609198
Height        3.128817
Diameter      -0.609198
Height        3.128817
Whole weight  0.530959
Shucked weight 0.719098
Viscera weight 0.591852
Shell weight  0.620927
Rings         1.114102
dtype: float64
```

```
In [120]: sns.boxplot(data['Rings'])
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-120-0f1ddeb65a1> in <module>
----> 1 sns.boxplot(data['Rings'])

TypeError: list indices must be integers or slices, not str
```

```
In [134]: q1=data['Rings'] . describe()['25%']
q3=data['Rings'] . describe()['75%']
q1
```

```
Out[134]: 8.0
```

Type Markdown and LaTeX: α^2

```
In [84]: q3
```

```
Out[84]: 11.0
```

```
In [85]: iqr=q3-q1
iqr
```

```
Out[85]: 3.0
```

```
In [86]: a=q1-(1.5*q1)
b=q3+(1.5*q3)
print(a, b)
```

```
-4.0 27.5
```

```
In [87]: data[data['Rings']<a]
```

```
Out[87]:   Sex  Length  Diameter  Height  Whole weight  Shucked weight  Viscera weight  Shell weight  Rings
```

```
In [88]: data[data['Rings']>b].head()
```


```
Out[88]:   Sex  Length  Diameter  Height  Whole weight  Shucked weight  Viscera weight  Shell weight  Rings
480  F      0.7    0.585    0.185    1.8075    0.7055    0.3215    0.475    29
```

```
In [89]: outlier_list=list(data[data['Rings']>b]['Rings'])
print(outlier_list)
```

```
[29]
```

```
In [90]: data['Rings']=data['Rings'].replace(outlier_list)
sns.boxplot(data['Rings'])
```

```
Out[90]: <AxesSubplot:xlabel='Rings'>
```



```
In [95]: le=LabelEncoder()
data['Sex']=le.fit_transform(data['Sex'])
data
```

```
Out[95]:   Sex  Length  Diameter  Height  Whole weight  Shucked weight  Viscera weight  Shell weight  Rings
0    2    0.455    0.365    0.095    0.5140    0.2245    0.1010    0.1500    15
1    2    0.350    0.265    0.090    0.2255    0.0995    0.0485    0.0700    7
2    0    0.530    0.420    0.135    0.6770    0.2565    0.1415    0.2100    9
3    2    0.440    0.365    0.125    0.5160    0.2155    0.1140    0.1550    10
4    1    0.330    0.255    0.080    0.2050    0.0895    0.0395    0.0550    7
...  ...  ...      ...      ...      ...      ...      ...      ...
4172  0    0.565    0.450    0.165    0.8870    0.3700    0.2390    0.2490    11
4173  2    0.590    0.440    0.135    0.9660    0.4390    0.2145    0.2605    10
```

4174	2	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	0	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	2	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12

4177 rows × 9 columns

```
In [96]: X = data.iloc[:, :-1].values
y = data.iloc[:, -1].values
```

```
In [97]: train_X, val_X, train_y, val_y = train_test_split(X, y, test_size = 0.2, random_state = 0)
print("Shape of Training X :", train_X.shape)
print("Shape of Validation X :", val_X.shape)
```

Shape of Training X : (3341, 8)
Shape of Validation X : (836, 8)

```
In [98]: lr = LinearRegression()
lr.fit(train_X, train_y)
print('Attempting to fit Linear Regressor')
```

Attempting to fit Linear Regressor

```
In [105]: %%time
y_pred_val_lr = lr.predict(val_X)
print('MAE on Validation set :', metrics.mean_absolute_error(val_y, y_pred_val_lr))
print("\n")
print('MSE on Validation set :', metrics.mean_squared_error(val_y, y_pred_val_lr))
print("\n")
print('RMSE on Validation set :', np.sqrt(metrics.mean_squared_error(val_y, y_pred_val_lr)))
print("\n")
print('R2 Score on Validation set :', metrics.r2_score(val_y, y_pred_val_lr))
print('R2 Score on Validation set :', metrics.r2_score(val_y, y_pred_val_lr))
print("\n")
```

MAE on Validation set : 1.5786845608962012

MSE on Validation set : 4.7449590677450635

RMSE on Validation set : 1.2564571464623062

R2 Score on Validation set : 0.5466388609280107

Wall time: 6 ms

In []: